

Status_Summary

HelixTranslate — Feature Status Summary

Revision: 10 **Last modified:** 2026-06-16T13:00:00Z **Authority:** Two-audience companion to docs/features/Status.md (§11.4.56). Derived from the same inventory. Per §11.4.44 (revision header), §11.4.60 (always-sync), §11.4.6 (no-guessing).

Page 1 — For the team (non-developer)

What is this? HelixTranslate is an ebook translation system. This page is a plain-language snapshot of how much of it we have catalogued and how much we have proven works on camera.

What works

- We have a complete catalogue of **every feature** in the product and its 8 owned add-on modules — **494 individual capabilities** catalogued (Rev 9 added 1 new unified-translator HTML→Markdown conversion row for the Wave-2 video; Rev 8 re-counted the live tables directly and added the 7 missing agent-bridge ensemble rows to reach 493; the headline “417” is the same list with closely-related sub-options counted once).
- The majority are fully built and reachable: the command-line translators, the web dashboards, the REST and gRPC servers, ~32 built-in AI providers, ebook format readers (FB2, EPUB, DOCX, HTML, TXT, PDF) and writers (EPUB, FB2, TXT, HTML, Markdown, **DOCX and now PDF**), caching, security, and all 8 add-on modules.
- **32 features are proven on real recordings** — each checked frame-by-frame to confirm it genuinely shows the feature working (real translated text in the right language, a live connection count changing, etc.), not just “a screen”. These include: the primary command-line translator doing real DeepSeek translations and converting between ebook formats (EPUB→TXT, HTML→EPUB, FB2→EPUB, FB2→FB2, **and now HTML→Markdown**), Serbian Cyrillic/Latin handling, the **multi-pass polishing engine (-multipass)**, the REST API server translating for real, the gRPC server translating English→Spanish, the markdown round-trip tool, the preparation/analysis runner (fixed), the simple

CLI translator, the live monitoring dashboard, the **agent bridge command-line tool**, the **agent bridge MCP-stdio server**, and — **new this round (Wave 2)** — **DOCX output**, **PDF output**, and the **provider-diverse translator set** (four real per-provider translations: novita, mistral, groq, siliconflow).

- **DOCX output and PDF output are now proven on camera (Wave 2).** A real run of the command-line translator produces a real .docx (“Microsoft Word 2007+” with real Spanish text) and a valid PDF with real Spanish — both captured as real-use videos and checked frame-by-frame. **The blocker is fixed:** the default command-line/gRPC translation path that was runtime-broken in Wave 1 (the auto-selected verified model being rejected by a stale model whitelist) now works — which is exactly why these recordings genuinely demonstrate working features. We never faked a “working” recording; in Wave 1 these were honest skips.
- **New this round (Wave 3a):** the **Cerebras** and **SambaNova** per-provider LLM clients are now video-proven — each a real EN→ES translation run through `unified-translator -provider <name>`, with the two Spanish outputs distinct from each other (“El sol de la mañana se alzó...” vs “El sol matutino se alzó...”), proving genuine, separate per-provider models rather than a canned response. (Prior round, Wave 2: DOCX output, PDF output, HTML→Markdown conversion, and the provider-diverse translators were video-proven; the Wave-1 default-translation-path blocker is resolved.)

What is pending or limited

- **Video proof is still the main gap: 30 of 494** features have a recorded, watch-it-yourself demonstration ($\approx 6.1\%$). Everything else is built and present in the code, but a recording is still owed. Nothing is claimed proven just because the code exists.
- **PDF output is now built** (this round). The product can *read* DOCX and PDF, and can now *write* DOCX **and** PDF. Output is now EPUB, FB2, TXT, HTML, Markdown, DOCX and PDF. (PDF writing needs the weasyprint tool, which is installed here.)
- **The web dashboard’s video proof is honestly skipped, not faked:** recording the dashboard translating end-to-end needs the HelixQA web/video test backend, which was off-limits to this round. The page is proven working by its tests + a real run, but its on-camera proof is owed.
- **The on-this-computer AI engines and the remote-worker translators were removed (this round):** the product no longer runs a local AI model (llama.cpp / Ollama) and no longer ships the SSH-based remote-worker tools (`translate-ssh`, `ssh-translation`, `ebook-translator`, the documented translator tool). Every translation now uses an online verified AI provider chosen automatically via the LLMsVerifier bridge. These removed tools are marked **Obsolete** (operator-approved 2026-06-15). The team-facing distributed/API capability is kept.
- **The “verified models” API returns 404** because the model-verification service is switched off in the current configuration (a config choice, not a broken feature).
- **4 features are “stubs”** — they look like real endpoints but return placeholder data: `ebook-translate-via-API`, the two preparation/analysis API endpoints, and `translate-with-verification`.
- A handful of small API endpoints return fixed/sample data (e.g. a stats endpoint that always shows zeros). These are flagged in the detailed document.

Video-confirmation coverage

32 / 494 ≈ 6.5% (Wave 3a added 2 net-new video confirmations: the Cerebras and SambaNova per-provider LLM clients, each a distinct real EN→ES translation; Wave 2 added 5: DOCX output, PDF output, HTML→Markdown conversion, ProviderDiverseTranslators, EnsembleFactory). This is the headline number to keep improving: each new genuine recording raises it.

Team actions

- **No operator decision owed for the former SSH/local tools:** the remote-worker + local-engine tools were removed (operator-approved 2026-06-15), so they no longer wait on a remote host — they are Obsolete, not blocked.
- Priority ask: produce real recordings for the next-most-important runtime features (web dashboard translation, the other AI providers) to keep lifting video-confirmation coverage.

Page 2 — For software engineers

Source & method

Derived 1:1 from docs/features/.feature_inventory_raw.md (Rev 1). Every detailed inventory row → one Status row. Statuses set per §11.4.6 from what the inventory actually found in source; no test files were attributed per-feature (recorded Not-inventoried, an honest unknown, never guessed).

Counts

Dimension	Value
Headline total (dedup per-category tally)	417
Enumerated detailed rows (this doc + Status.md)	494 (Rev 9 = 493 Rev 8 + 1 new unified-translator HTML→Markdown conversion row added for the Wave-2 html-to-md video; Rev 8 added the 7 missing pkg/bridge ensemble-seam rows [BestTranslator, BestTranslatorFunc, EnsembleFactory, ProviderDiverseTranslators, ProviderDiverseClients, ProviderDiverseModels, BestClient] → 493)
Implemented	

Dimension	Value
	443 (Rev 9: 442 Rev 8 + 1 new unified-translator HTML→Markdown conversion row; Rev 8: 435 Rev 7 + 7 new pkg/bridge ensemble-seam rows; was 471 Rev 6, -36 flipped to Obsolete Rev 7)
Obsolete (→ Fixed.md)	39 (Rev 7 — bridge phase-2 R-2..R-4 removals: cmd/translate-ssh/ssh-translation/translator/ebook-translator, pkg/sshworker, pkg/modelsbridge, Ollama + llama.cpp providers; Reason=feature-removed)
Stub	4
Not implemented (gap rows)	0 (PDF-write flipped gap→Implemented Rev 4, commit fb265e7)
Partial	2 (POST /api/batch; vision_engine OpenCV) — was 3 Rev 6; cmd/translator → Obsolete Rev 7
Operator-blocked	0 (was 2 Rev 6; translate-ssh + ebook-translator → Obsolete Rev 7 — capability removed, not host-blocked)
Video-confirmed	32 (Wave 3a added 2 net-new: Cerebras + SambaNova per-provider LLM clients, distinct real EN→ES translations; Wave 2 added 5: DOCX output, PDF output, HTML→Markdown conversion, ProviderDiverseTranslators, EnsembleFactory)
Video PENDING	runtime/user-visible, no recording yet
Video N/A	flags / internal types / infra middleware

Per-category (inventory headline)

Category	Count
CLI	130
API	70
Web	23
gRPC	7
Library	81
Submodule	95
Infra	11
TOTAL	417

Stub list (Implementation=Stub — returns mock, no real work)

1. POST /api/v1/translate/ebook — pkg/api/handler.go:1954 (emits start event, returns mock; does NOT translate).
2. POST /api/v1/preparation/analyze — pkg/api/handler.go:1832 (file stat/count + events; stub analysis, no LLM).
3. GET /api/v1/preparation/result/:session_id — pkg/api/handler.go:1927 (returns mock result).
4. POST /api/v1/translate-with-verification — pkg/api/verifier_handlers.go (selects/verifies a model, returns metadata + preview; does NOT translate).

Partial list

1. POST /api/batch (standalone pkg/api/server.go) — returns queued batch_id only, no translation executed.
2. vision_engine OpenCV — stub by default; real CV only behind a build tag.

cmd/translator was Partial in Rev 6 (local STUB + remote OPERATOR-BLOCKED) — flipped to **Obsolete** Rev 7 (the whole binary was removed in bridge phase-2 R-4).

Not-implemented gap rows (§11.4.6)

None — both former gap rows are now Implemented.

Resolved Rev 3: pkg/ebook DOCX-write — Implemented (pure-Go OOXML writer pkg/ebook/docx_writer.go, commit 87cd2be; real DeepSeek run → garden_es.docx = Microsoft Word 2007+). Resolved Rev 4: pkg/ebook PDF-write — Implemented (pkg/ebook/pdf_writer.go, commit fb265e7; Book→HTML5→weasyprint→Cyrillic-faithful PDF, wired into unified-translator .pdf output, main.go:1160; pkg/ebook PDF tests green). Output formats are now EPUB / FB2 / TXT / HTML / MD / DOCX / PDF.

New Rev 4 features (built + tested this round, recordings owed/skipped)

1. **cmd/model-bridge + pkg/bridge** — LLMsVerifier→component+agent bridge (no local models): selects the best verifier-scored model (top-1 + fallback) and exposes it via a CLI and an MCP stdio server (.mcp.json). Commits ab1bed3 + a5860b2. PASS (real run: best-model novita/Sao10K/L3-8B-Stheno-v3.2 score 0.919, Invoke→“Buenos días, amigo.”; MCP nonce-echo HELIX-PR00F-9b21x; go test -race ./pkg/bridge green; §1.1 routing-guard TestBridge_Invoke_RoutesToBestModel). **CLI video CONFIRMED** (helixtranslate-bridge-bestmodel-translate-20260615.mp4, ffprobe 8.0s/80fr, §11.4.107 content-verified); **MCP-stdio surface video**

CONFIRMED (helixtranslate-bridge-mcp-stdio-20260615.mp4 — real JSON-RPC MCP session: initialize→tools/list→bridge_best_model→bridge_invoke, tool bridge_invoke real output “Le pont relie deux rives.” isError=None; fprobe 11.88s/10fr, ≥6 frames §11.4.107 content-verified).

2. **internal/verifier affirmative-response hard-gate** — commit 97a8afd: a model with affirmative_response=0 is now a hard disqualifier. PASS (§11.4.115 polarity guard pipeline_affirmative_gate_test.go: GREEN-default — the standing suite asserts the defect ABSENT; RED reproduction is OPT-IN only under RED_MODE=1, green this session). Video N/A (internal gate, no user-visible surface).
3. **pkg/api/dashboard.go web-dashboard backend** — commit f6ba5cc: dashboard page + translations endpoints, previously 404, now genuinely translates. PASS unit/integration (dashboard_test.go §11.4.115 RED→GREEN, go test -race ./pkg/api green; live run → “Dobro jutro, prijatelju.”). Video **SKIP** (§11.4.3/§11.4.52 — autonomous browser+video capture needs the HelixQA web/video backend, off-limits this round; migration item owed; NOT a faked confirmation).

Operator-blocked list (§11.4.45)

None — the Rev 6 operator-blocked binaries (cmd/translate-ssh, cmd/ebook-translator) were **removed** in bridge phase-2 R-4 and are now Obsolete (see the Obsolete list), not host-blocked.

Obsolete list (§11.4.90 — Rev 7, operator-approved 2026-06-15 D1/D2 + R-1 + R-4)

1. cmd/translate-ssh, cmd/ssh-translation, cmd/translator, cmd/ebook-translator — SSH-local / remote-worker translators (removed; FACT: git ls-files empty for these cmd dirs).
2. pkg/sshworker, pkg/modelsbridge — SSH worker library + models-bridge adapter (removed).
3. Ollama provider, llama.cpp provider + multi-worker coordinator (pkg/translator/llm/{ollama.go, llamacpp.go, llamacpp_provider.go} removed).

Reason=feature-removed; Superseding-item=pkg/bridge (LLMsVerifier→agent bridge); the KEPT distributed-API path is NOT Obsolete. Triple-check: source files gone + CM-NO-LOCAL-RUNTIME gate PASS.

Static / hardcoded-return endpoints (Implemented, but not live)

- GET /api/stats (pkg/api/server.go) — hardcoded zeros.
- GET /api/v1/status/:session_id (pkg/api/handler.go) — hardcoded completed.
- GET /api/v1/metrics (cmd/api-server/main.go) — static/zero.

- GET /api/v1/status/:session_id (cmd/monitor-server/main.go) — static monitoring_active.
- GET /api/v1/providers (gRPC GetProviders) — static ProviderRegistry (API providers only; the ssh local provider was removed in bridge phase-2 R-4).
- Static language/provider lists in several GET /api/.../languages routes.

Top gaps (engineering)

1. **Video-confirmation 32/494.** 32 rows carry real, ffmpeg- and content-verified recordings. **Wave 3a (2026-06-16) added 2 net-new** (validator-style ffmpeg+frame, §11.4.107): the **Cerebras** client (helix_translate-provider-cerebras-translate-wave3a-20260616.mp4, “El sol de la mañana se alzó sobre el tranquilo pueblo.”) and the **SambaNova** client (helix_translate-provider-sambanova-translate-wave3a-20260616.mp4, “El sol matutino se alzó sobre el tranquilo pueblo.”), both via unified-translator -provider <name>, the two outputs distinct from each other (genuine per-provider models), start-frame ≠ output-frame liveness. **Wave 2 (2026-06-16) added 5 net-new** (validator-verified ffmpeg+frame, §11.4.107): DOCX output (helix_translate-docx-output-20260616-112943.mp4), PDF output (helix_translate-pdf-output-20260616-113057.mp4), HTML→Markdown conversion (helix_translate-html-to-md-20260616-114228.mp4), and ProviderDiverseTranslators + EnsembleFactory (the 4 per-provider clips helix_translate-provider-{novita,mistral,groq,siliconflow}-translate-20260616-114*.mp4); plus an additional re-confirming clip on the already-confirmed gRPC TranslationService row (helix_translate-grpc-translate-20260616-113415.mp4, not double-counted). Wave 1 added 2 (bridge ensemble + provider-diverse catalogue). Prior recordings: -multipass, bridge CLI, bridge MCP-stdio, the 3 video-surfaced bug fixes in commit a5e8866. Per §11.4.2/§11.4.107 nothing may be marked video-confirmed without a real, content-verified file.
 -  **Wave-1 DISCOVERED BLOCKER — RESOLVED (Wave 2, §11.4.108/§11.4.138):** the default bridge **translator** path (unified-translator + gRPC) that was RUNTIME-BROKEN in Wave 1 (VerifiedFactory.CreateTranslatorWithFallback→NewLLMTranslatorWithConfig rejecting the verified strongest model against the stale ValidModels whitelist, llm.go:237) is now FIXED — the Wave-2 DOCX/PDF/gRPC/ per-provider recordings are real-use videos showing the verified-bridge translator path producing real target-language output (the exact path that failed in Wave 1). The honest Wave-1 SKIPS are now genuine CONFIRMED rows. See Status.md caveat 8 for the preserved forensic record.
2. **PDF output now Implemented** (Rev 4, commit fb265e7) — pkg/ebook now has FB2/EPUB/DOCX/PDF writers; zero open gap rows remain. Output is EPUB/FB2/TXT/HTML/MD/DOCX/PDF.
3. **-multipass wired (d53e085); -verify ≠ -multipass** — unified-translator -verify runs the CLI’s own per-step check; the new -multipass flag invokes the pkg/verification multi-pass polisher engine (now wired + video-confirmed). cmd/cli write-safety FIXED (87cd2be, atomic write/no-clobber, guard cmd/cli/no_partial_output_test.go) + title path-leak FIXED (d53e085).
4. **SSH-local + local-runtime features removed (Rev 7, 39 rows Obsolete)** — translate-ssh/ssh-translation/translator/ebook-translator binaries,

pkg/sshworker, pkg/modelsbridge, and the Ollama + llama.cpp providers are gone (operator-approved); the default path now sources the LLMsVerifier bridge.

5. **GET /api/v1/verified-models** → **404** — LLMsVerifier disabled in config (route not mounted); feature-disabled-by-config, not a defect.
6. **Per-feature test attribution missing** (Not-inventoried everywhere). Repo has extensive `_test.go` + `test/` suites; mapping suites → features is future work.
7. **4 stubs + 1 unimplemented + 3 partials** must either be implemented or reclassified per §11.4.90 before they can be claimed as working capabilities.
8. **UNCONFIRMED LLM determinations** carried from the inventory: thin OpenAI-compatible providers' base-URL/model specifics; several `llms_verifier` packages (scoring/helixqa/cliagents/crush/opencode/scheduler/partners/bigdata/multimodal/performance); challenges/userflow evaluators; `doc_processor/archdoc`.

Video-coverage ratio

32 / 494 enumerated rows = **6.5%** ($\approx 7.7\%$ against the 417 headline). Wave 3a added 2 net-new video confirmations (Cerebras + SambaNova per-provider clients, distinct real EN→ES translations); Wave 2 added 5 (DOCX/PDF/HTML→MD/ProviderDiverseTranslators/EnsembleFactory). This is the primary metric to drive upward; each anti-bluff §11.4.107 recording increments the numerator.

Cross-references

- Detail: `docs/features/Status.md`
- Source: `docs/features/.feature_inventory_raw.md`
- Sync context: `.docs_chain/contexts/features.yaml` (§11.4.106)