

Status_Summary

HelixTranslate — Feature Status Summary

Revision: 7 **Last modified:** 2026-06-16T00:00:00Z **Authority:** Two-audience companion to docs/features/Status.md (§11.4.56). Derived from the same inventory. Per §11.4.44 (revision header), §11.4.60 (always-sync), §11.4.6 (no-guessing).

Page 1 — For the team (non-developer)

What is this? HelixTranslate is an ebook translation system. This page is a plain-language snapshot of how much of it we have catalogued and how much we have proven works on camera.

What works

- We have a complete catalogue of **every feature** in the product and its 8 owned add-on modules — **478 individual capabilities** catalogued (470 from the source inventory + 2 output rows [DOCX + PDF output, both now built] + 6 new rows this round [the agent bridge CLI+MCP, the bridge library, the model-verification fix, the web-dashboard backend]; the headline “417” is the same list with closely-related sub-options counted once).
- The majority are fully built and reachable: the command-line translators, the web dashboards, the REST and gRPC servers, ~32 built-in AI providers, ebook format readers (FB2, EPUB, DOCX, HTML, TXT, PDF) and writers (EPUB, FB2, TXT, HTML, Markdown, **DOCX and now PDF**), caching, security, and all 8 add-on modules.
- **23 features are proven on real recordings** — each checked frame-by-frame to confirm it genuinely shows the feature working (real translated text in the right language, a live connection count changing, etc.), not just “a screen”. These include: the primary command-line translator doing real DeepSeek translations and converting between ebook formats (EPUB→TXT, HTML→EPUB, FB2→EPUB, FB2→FB2), Serbian Cyrillic/Latin handling, the **multi-pass polishing engine (-multipass)**, the REST API server translating for real, the gRPC server translating English→Spanish, the markdown round-trip tool, the preparation/analysis runner (fixed), the simple CLI translator, the live

monitoring dashboard, **the agent bridge command-line tool** (picks the best verifier-scored model and translates through it — recording shows it choosing novita/Sao10K/L3-8B-Stheno-v3.2 and translating “Good morning, friend.” → “Buenos días, amigo.”), and **the agent bridge MCP-stdio server** (a real JSON-RPC MCP session — recording shows tools/list, bridge_best_model, and the bridge_invoke tool returning a live model translation “Le pont relie deux rives.”).

- **DOCX output** is built and proven by the produced file itself (a real .docx = “Microsoft Word 2007+” with real translated text). **PDF output is newly built this round** and proven by passing tests (it renders the translated book to a real, valid PDF that keeps Serbian Cyrillic readable). Screen recordings of both are still owed, so they count as file/test-proven rather than video-proven for now.
- **New this round:** the **agent bridge MCP-stdio server** is now **video-proven** (helixtranslate-bridge-mcp-stdio-20260615.mp4 — a real JSON-RPC MCP session over stdio: initialize→tools/list→bridge_best_model→bridge_invoke, with the bridge_invoke tool returning a live verified-model translation “Le pont relie deux rives.”), joining the **agent bridge command-line tool** which was video-proven the previous round (helixtranslate-bridge-bestmodel-translate-20260615.mp4); a **correctness fix in the model-verification gate** (a model that fails the basic “do you respond?” check is now properly rejected); and the **web dashboard’s translation page** which used to return “page not found” and now genuinely translates.

What is pending or limited

- **Video proof is still the main gap: 23 of 478** features have a recorded, watch-it-yourself demonstration (≈ 4.8%). Everything else is built and present in the code, but a recording is still owed. Nothing is claimed proven just because the code exists.
- **PDF output is now built** (this round). The product can *read* DOCX and PDF, and can now *write* DOCX and PDF. Output is now EPUB, FB2, TXT, HTML, Markdown, DOCX and PDF. (PDF writing needs the weasyprint tool, which is installed here.)
- **The web dashboard’s video proof is honestly skipped, not faked:** recording the dashboard translating end-to-end needs the HelixQA web/video test backend, which was off-limits to this round. The page is proven working by its tests + a real run, but its on-camera proof is owed.
- **The on-this-computer AI engines and the remote-worker translators were removed (this round):** the product no longer runs a local AI model (llama.cpp / Ollama) and no longer ships the SSH-based remote-worker tools (translate-ssh, ssh-translation, ebook-translator, the documented translator tool). Every translation now uses an online verified AI provider chosen automatically via the LLMsVerifier bridge. These removed tools are marked **Obsolete** (operator-approved 2026-06-15). The team-facing distributed/API capability is kept.
- **The “verified models” API returns 404** because the model-verification service is switched off in the current configuration (a config choice, not a broken feature).
- **4 features are “stubs”** — they look like real endpoints but return placeholder data: ebook-translate-via-API, the two preparation/analysis API endpoints, and translate-with-verification.

- A handful of small API endpoints return fixed/sample data (e.g. a stats endpoint that always shows zeros). These are flagged in the detailed document.

Video-confirmation coverage

23 / 478 $\approx$ 4.8%, plus **2 real-artifact/test-confirmed** (DOCX + PDF output, video PENDING). This is the headline number to keep improving: each new genuine recording raises it.

Team actions

- **No operator decision owed for the former SSH/local tools:** the remote-worker + local-engine tools were removed (operator-approved 2026-06-15), so they no longer wait on a remote host — they are Obsolete, not blocked.
- Priority ask: produce real recordings for the next-most-important runtime features (web dashboard translation, the other AI providers) to keep lifting video-confirmation coverage.

Page 2 — For software engineers

Source & method

Derived 1:1 from docs/features/.feature_inventory_raw.md (Rev 1). Every detailed inventory row → one Status row. Statuses set per §11.4.6 from what the inventory actually found in source; no test files were attributed per-feature (recorded Not-inventoried, an honest unknown, never guessed).

Counts

Dimension	Value
Headline total (dedup per-category tally)	417
Enumerated detailed rows (this doc + Status.md)	478 (470 inventory + 2 output rows [DOCX-write Implemented Rev 3, PDF-write Implemented Rev 4] + 6 new Rev 4 rows: cmd/model-bridge ×2, pkg/bridge ×2, internal/verifier gate ×1, pkg/api/dashboard.go ×1)
Implemented	435 (was 471 Rev 6; −36 flipped to Obsolete Rev 7)
Obsolete (→ Fixed.md)	39 (Rev 7 — bridge phase-2 R-2..R-4 removals: cmd/translate-ssh/ssh-translation/translator/ebook-

Dimension	Value
Stub	translator, pkg/sshworker, pkg/modelsbridge, Ollama + llama.cpp providers; Reason=feature-removed) 4
Not implemented (gap rows)	0 (PDF-write flipped gap→Implemented Rev 4, commit fb265e7)
Partial	2 (POST /api/batch; vision_engine OpenCV) — was 3 Rev 6; cmd/translator → Obsolete Rev 7
Operator-blocked	0 (was 2 Rev 6; translate-ssh + ebook-translator → Obsolete Rev 7 — capability removed, not host-blocked)
Video-confirmed	23 (+ 2 real-artifact/test-confirmed: DOCX + PDF output)
Video PENDING	runtime/user-visible, no recording yet
Video N/A	flags / internal types / infra middleware

Per-category (inventory headline)

Category	Count
CLI	130
API	70
Web	23
gRPC	7
Library	81
Submodule	95
Infra	11
TOTAL	417

Stub list (Implementation=Stub — returns mock, no real work)

1. POST /api/v1/translate/ebook — pkg/api/handler.go:1954 (emits start event, returns mock; does NOT translate).
2. POST /api/v1/preparation/analyze — pkg/api/handler.go:1832 (file stat/count + events; stub analysis, no LLM).
3. GET /api/v1/preparation/result/:session_id — pkg/api/handler.go:1927 (returns mock result).
4. POST /api/v1/translate-with-verification — pkg/api/verifier_handlers.go (selects/verifies a model, returns metadata + preview; does NOT translate).

Partial list

1. `POST /api/batch` (standalone `pkg/api/server.go`) — returns queued `batch_id` only, no translation executed.
2. `vision_engine` OpenCV — stub by default; real CV only behind a build tag.

`cmd/translator` was Partial in Rev 6 (local STUB + remote OPERATOR-BLOCKED) — flipped to **Obsolete** Rev 7 (the whole binary was removed in bridge phase-2 R-4).

Not-implemented gap rows (§11.4.6)

None — both former gap rows are now Implemented.

Resolved Rev 3: `pkg/ebook` DOCX-write — Implemented (pure-Go OOXML writer `pkg/ebook/docx_writer.go`, commit 87cd2be; real DeepSeek run → `garden_es.docx` = Microsoft Word 2007+). Resolved Rev 4: `pkg/ebook` PDF-write — Implemented (`pkg/ebook/pdf_writer.go`, commit fb265e7; Book→HTML5→weasyprint→Cyrillic-faithful PDF, wired into unified-translator .pdf output, main.go:1160; `pkg/ebook` PDF tests green). Output formats are now EPUB / FB2 / TXT / HTML / MD / DOCX / PDF.

New Rev 4 features (built + tested this round, recordings owed/skipped)

1. **`cmd/model-bridge` + `pkg/bridge`** — LLMsVerifier→component+agent bridge (no local models): selects the best verifier-scored model (top-1 + fallback) and exposes it via a CLI and an MCP stdio server (`.mcp.json`). Commits ab1bed3 + a5860b2. PASS (real run: best-model novita/Sao10K/L3-8B-Stheno-v3.2 score 0.919, Invoke→“Buenos días, amigo.”; MCP nonce-echo HELIX-PR00F-9b21x; `go test -race ./pkg/bridge` green; §1.1 routing-guard TestBridge_Invoke_RoutesToBestModel). **CLI video CONFIRMED** (`helixtranslate-bridge-bestmodel-translate-20260615.mp4`, ffprobe 8.0s/80fr, §11.4.107 content-verified); **MCP-stdio surface video CONFIRMED** (`helixtranslate-bridge-mcp-stdio-20260615.mp4` — real JSON-RPC MCP session: initialize→tools/list→bridge_best_model→bridge_invoke, tool bridge_invoke real output “Le pont relie deux rives.” isError:None; ffprobe 11.88s/10fr, ≥6 frames §11.4.107 content-verified).
2. **`internal/verifier affirmative-response hard-gate`** — commit 97a8afd: a model with `affirmative_response=0` is now a hard disqualifier. PASS (§11.4.115 polarity guard pipeline `affirmative_gate_test.go`: GREEN-default — the standing suite asserts the defect ABSENT; RED reproduction is OPT-IN only under `RED_MODE=1`, green this session). Video N/A (internal gate, no user-visible surface).
3. **`pkg/api/dashboard.go web-dashboard backend`** — commit f6ba5cc: dashboard page + translations endpoints, previously 404, now genuinely translates. PASS unit/integration (`dashboard_test.go` §11.4.115 RED→GREEN, `go test -race ./pkg/api` green; live run → “Dobro jutro,

prijateljju.”). Video **SKIP** (§11.4.3/§11.4.52 — autonomous browser+video capture needs the HelixQA web/video backend, off-limits this round; migration item owed; NOT a faked confirmation).

Operator-blocked list (§11.4.45)

None — the Rev 6 operator-blocked binaries (cmd/translate-ssh, cmd/ebook-translator) were **removed** in bridge phase-2 R-4 and are now Obsolete (see the Obsolete list), not host-blocked.

Obsolete list (§11.4.90 — Rev 7, operator-approved 2026-06-15 D1/D2 + R-1 + R-4)

1. cmd/translate-ssh, cmd/ssh-translation, cmd/translator, cmd/ebook-translator — SSH-local / remote-worker translators (removed; FACT: git ls-files empty for these cmd dirs).
2. pkg/sshworker, pkg/modelsbridge — SSH worker library + models-bridge adapter (removed).
3. Ollama provider, llama.cpp provider + multi-worker coordinator (pkg/translator/llm/{ollama.go, llamacpp.go, llamacpp_provider.go} removed).

Reason=feature-removed; Superseding-item=pkg/bridge (LLMsVerifier→agent bridge); the KEPT distributed-API path is NOT Obsolete. Triple-check: source files gone + CM-NO-LOCAL-RUNTIME gate PASS.

Static / hardcoded-return endpoints (Implemented, but not live)

- GET /api/stats (pkg/api/server.go) — hardcoded zeros.
- GET /api/v1/status/:session_id (pkg/api/handler.go) — hardcoded completed.
- GET /api/v1/metrics (cmd/api-server/main.go) — static/zero.
- GET /api/v1/status/:session_id (cmd/monitor-server/main.go) — static monitoring_active.
- GET /api/v1/providers (gRPC GetProviders) — static ProviderRegistry (API providers only; the ssh local provider was removed in bridge phase-2 R-4).
- Static language/provider lists in several GET /api/.../languages routes.

Top gaps (engineering)

1. **Video-confirmation 23/478** (+ 2 real-artifact/test-confirmed: DOCX + PDF output). 23 rows carry real, ffprobe- and content-verified recordings (Status.md cites each .mp4; the -multipass video is helixtranslate-cli-multipass-verify-20260615.mp4, ffprobe 888x630/93fr/9.3s; the bridge CLI video is

- helixtranslate-bridge-bestmodel-translate-20260615.mp4, ffprobe 790×560/80fr/8.0s; the bridge MCP-stdio video is helixtranslate-bridge-mcp-stdio-20260615.mp4, ffprobe 790×560/10fr/11.88s; the 3 video-surfaced bug fixes — prep-translator dead, server TLS startup, REST hardcoded target-lang — are in commit a5e8866). The remaining Rev-4 additions carry PENDING (PDF), SKIP (web dashboard — needs HelixQA web/video backend) or N/A (internal verifier gate). Per §11.4.2/§11.4.107 nothing may be marked video-confirmed without a real, content-verified file.
2. **PDF output now Implemented** (Rev 4, commit fb265e7) — pkg/ebook now has FB2/EPUB/DOCX/PDF writers; zero open gap rows remain. Output is EPUB/FB2/TXT/HTML/MD/DOCX/PDF.
 3. **-multipass wired (d53e085); -verify ≠ -multipass** — unified-translator -verify runs the CLI's own per-step check; the new -multipass flag invokes the pkg/verification multi-pass polisher engine (now wired + video-confirmed). cmd/cli write-safety FIXED (87cd2be, atomic write/no-clobber, guard cmd/cli/no_partial_output_test.go) + title path-leak FIXED (d53e085).
 4. **SSH-local + local-runtime features removed (Rev 7, 39 rows Obsolete)** — translate-ssh/ssh-translation/translator/ebook-translator binaries, pkg/sshworker, pkg/modelsbridge, and the Ollama + llama.cpp providers are gone (operator-approved); the default path now sources the LLMsVerifier bridge.
 5. **GET /api/v1/verified-models → 404** — LLMsVerifier disabled in config (route not mounted); feature-disabled-by-config, not a defect.
 6. **Per-feature test attribution missing** (Not-inventoried everywhere). Repo has extensive _test.go + test/ suites; mapping suites → features is future work.
 7. **4 stubs + 1 unimplemented + 3 partials** must either be implemented or reclassified per §11.4.90 before they can be claimed as working capabilities.
 8. **UNCONFIRMED LLM determinations** carried from the inventory: thin OpenAI-compatible providers' base-URL/model specifics; several llms_verifier packages (scoring/helixqa/cliagents/crush/opencode/scheduler/partners/bigdata/multimodal/performance); challenges/userflow evaluators; doc_processor/archdoc.

Video-coverage ratio

23 / 478 enumerated rows = **4.8%** ($\approx 5.5\%$ against the 417 headline). Plus 2 real-artifact/test-confirmed (DOCX + PDF output, video PENDING). This is the primary metric to drive upward; each anti-bluff §11.4.107 recording increments the numerator.

Cross-references

- Detail: docs/features/Status.md
- Source: docs/features/.feature_inventory_raw.md
- Sync context: .docs_chain/contexts/features.yaml (§11.4.106)