

Status_Summary

HelixTranslate — Feature Status Summary

Revision: 1 **Last modified:** 2026-06-15T13:44:25Z **Authority:** Two-audience companion to docs/features/Status.md (§11.4.56). Derived from the same inventory. Per §11.4.44 (revision header), §11.4.60 (always-sync), §11.4.6 (no-guessing).

Page 1 — For the team (non-developer)

What is this? HelixTranslate is an ebook translation system. This page is a plain-language snapshot of how much of it we have catalogued and how much we have proven works on camera.

What works

- We have a complete catalogue of **every feature** in the product and its 8 owned add-on modules — **470 individual capabilities** (the headline “417” is the same list with closely-related sub-options counted once).
- The vast majority — **464 of 470** — are fully built and reachable in the product: the command-line translators, the web dashboards, the REST and gRPC servers, ~32 built-in AI providers, ebook format readers/writers (FB2, EPUB, DOCX, HTML, TXT, PDF), caching, security, and all 8 add-on modules.
- **One feature is proven on a real recording:** the primary command-line translator (`unified-translator`) doing a real English → Spanish translation with the DeepSeek AI provider. The fix it depended on was verified in that recording.

What is pending or limited

- **Video proof is the big gap.** Only **1 of 470** features has a recorded, watch-it-yourself demonstration so far (about 0.2%). Everything else is built and present in the code, but a recording proving it end-to-end is still owed. This is shown honestly — we are not claiming a feature is proven just because the code exists.

- **4 features are “stubs”** — they look like real endpoints but currently return placeholder data instead of doing the real work: ebook-translate-via-API, the two preparation/analysis API endpoints, and translate-with-verification.
- **2 features are partial:** one batch API endpoint just queues without translating, and the image-analysis (OpenCV) module ships a placeholder unless a special build option is turned on.
- A handful of small API endpoints return fixed/sample data (e.g. a stats endpoint that always shows zeros). These are flagged in the detailed document.

Video-confirmation coverage

1 / 470 $\approx$ 0.2%. This is the headline number to improve: each future recording raises it.

Team actions

- No operator decision is currently blocked.
 - Priority ask: produce real recordings for the next-most-important runtime features (web dashboard translation, gRPC translation, the other AI providers) to lift the video-confirmation coverage.
-

Page 2 — For software engineers

Source & method

Derived 1:1 from docs/features/.feature_inventory_raw.md (Rev 1). Every detailed inventory row $\rightarrow$ one Status row. Statuses set per §11.4.6 from what the inventory actually found in source; no test files were attributed per-feature (recorded Not-inventoried, an honest unknown, never guessed).

Counts

Dimension	Value
Headline total (dedup per-category tally)	417
Enumerated detailed rows (this doc + Status.md)	470
Implemented	464
Stub	4
Partial	2
Video-confirmed	1
Video PENDING	runtime/user-visible, no recording yet

Dimension	Value
Video N/A	flags / internal types / infra middleware

Per-category (inventory headline)

Category	Count
CLI	130
API	70
Web	23
gRPC	7
Library	81
Submodule	95
Infra	11
TOTAL	417

Stub list (Implementation=Stub — returns mock, no real work)

1. POST /api/v1/translate/ebook — pkg/api/handler.go:1954 (emits start event, returns mock; does NOT translate).
2. POST /api/v1/preparation/analyze — pkg/api/handler.go:1832 (file stat/count + events; stub analysis, no LLM).
3. GET /api/v1/preparation/result/:session_id — pkg/api/handler.go:1927 (returns mock result).
4. POST /api/v1/translate-with-verification — pkg/api/verifier_handlers.go (selects/verifies a model, returns metadata + preview; does NOT translate).

Partial list

1. POST /api/batch (standalone pkg/api/server.go) — returns queued batch_id only, no translation executed.
2. vision_engine OpenCV — stub by default; real CV only behind a build tag.

Static / hardcoded-return endpoints (Implemented, but not live)

- GET /api/stats (pkg/api/server.go) — hardcoded zeros.
- GET /api/v1/status/:session_id (pkg/api/handler.go) — hardcoded completed.
- GET /api/v1/metrics (cmd/api-server/main.go) — static/zero.
- GET /api/v1/status/:session_id (cmd/monitor-server/main.go) — static monitoring_active.

- GET /api/v1/providers (gRPC GetProviders) — static ProviderRegistry (openai, anthropic, ssh).
- Static language/provider lists in several GET /api/.../languages routes.

Top gaps (engineering)

1. **Video-confirmation 1/470.** The single recorded proof is the unified-translator DeepSeek EN→Spanish run: /Volumes/T7/Downloads/Recordings/helixtranslate-cli-deepseek-translation-FIXED_20260615_163824.mp4. All other runtime features carry PENDING. Per §11.4.2/§11.4.107 nothing else may be marked video-confirmed without a real file.
2. **Per-feature test attribution missing** (Not-inventoried everywhere). Repo has extensive _test.go + test/ suites; mapping suites → features is future work.
3. **4 stubs + 2 partials** must either be implemented or reclassified per §11.4.90 before they can be claimed as working capabilities.
4. **UNCONFIRMED LLM determinations** carried from the inventory: thin OpenAI-compatible providers' base-URL/model specifics; several llms_verifier packages (scoring/helixqa/cliagents/crush/opencode/scheduler/partners/bigdata/multimodal/performance); challenges/userflow evaluators; doc_processor/archdoc.

Video-coverage ratio

1 / 470 enumerated rows = **0.21%** ($\approx 0.24\%$ against the 417 headline). This is the primary metric to drive upward; each anti-bluff §11.4.107 recording increments the numerator.

Cross-references

- Detail: docs/features/Status.md
- Source: docs/features/.feature_inventory_raw.md
- Sync context: .docs_chain/contexts/features.yaml (§11.4.106)