

README

Fix — stale provider model-allowlists (gemini/groq/zhipu) + live-proof (§11.4.107/§11.4.6)

Revision: 1 **Last modified:** 2026-06-14T14:20:00Z

The central ValidModels allowlist (pkg/translator/llm/llm.go) only accepted DEPRECATED models for gemini/groq/zhipu and REJECTED current ones at the LOCAL gate, so those providers could not even attempt a current model (the request never reached the API). Added current models (additive — old names kept for back-compat).

RED → GREEN (gate level, live-proven)

Before: model 'X' is not valid for provider 'Y'. Valid models: [<stale list>] (local gate, never hit the API). After: the current models PASS the gate and REACH the live provider API:

Provider	Model added	Live API response (proves it reached the provider)	Verdict
groq	llama-3.3-70b-versatile	HTTP 429 "Rate limit reached for model llama-3.3-70b-versatile"	model VALID (recognized; TPD quota exhausted today)
gemini	gemini-1.5-flash (+1.5-pro/2.0/2.5)	HTTP 400 "API Key not found"	model+gate OK; GEMINI_API_KEY invalid/expired (operator)
zhipu	glm-4-flash (+plus/air/...)	HTTP 400 code 1211 "模型不存在"	gate OK; this key's account/endpoint lacks the model (operator/needs zipu-client investigation)

Honest boundary (§11.4.6)

This fix removes the stale-allowlist BLOCKER and is live-proven to let current models reach the providers. FULL end-to-end translation for gemini/zhpu still depends on valid credentials/quota/account (operator-side) — NOT claimed as working here. DeepSeek and Mistral already prove full end-to-end translation (see `e2e_format_matrix_*` / `e2e_multiprovider_*`). A separate follow-up may be warranted: gemini/zhpu CLIENTS may have an endpoint/model-format issue like the Qwen one fixed earlier (the live APIs reject documented models) — flagged, not guessed.