

helixagent-comprehensive-analysis

COMPREHENSIVE ANALYSIS: HelixAgent Repository with Special Focus on LLMsVerifier Integration

Repository: <https://github.com/HelixDevelopment/HelixAgent> **Module:** `dev.helix.agent` **Language:** Go 1.26 **Analysis Date:** 2026 **Method:** GitHub API + Raw Content Endpoints

A. PROJECT OVERVIEW


HelixAgent is a **production-ready, AI-powered ensemble LLM service** written in Go that:

- Aggregates responses from **51+ LLM providers** to provide the most accurate outputs
- Exposes an **OpenAI-compatible REST API** + gRPC facade
- Implements **multi-round AI debate orchestration** (5 positions x 5 LLMs = 25 total)
- Supports **MCP (Model Context Protocol)**, **ACP (Agent Coordination Protocol)**, **LSP (Language Server Protocol)**
- Provides **embeddings, vision, RAG, and containerized infrastructure**
- Uses **dynamic provider selection** via LLMsVerifier verification scores
- Monorepo with **~60 submodules** across 8 phases of extraction

Key Metrics

- **51+ LLM providers** supported (see `internal/llm/providers/`)
- **13 embedding providers**
- **35 MCP implementations**
- **10 LSP servers**
- **24+ power features**
- **65.6% test coverage**
- **193+ validation challenge scripts** with 1500+ tests
- **48 CLI agents** supported with auto-generated configs
- **32+ code formatters** (11 native, 14 service, 7 built-in)
- **45+ MCP adapters**

Architecture Diagram



B. COMPLETE DIRECTORY STRUCTURE

HelixAgent/	
— cmd/	# 8 application entry points
— helixagent/	# Main production server (port 8100)
— api/	# Standalone demo API (port 8080)
— grpc-server/	# gRPC service endpoint
— mcp-bridge/	# MCP SSE bridge (port 8103)
— cognee-mock/	# Mock Cognee service
— sanity-check/	# System validation tool
— generate-constitution/	# Constitution generator
— audit/	# Audit utility
— internal/	# Core application code (~50+ packages)
— adapters/	# External submodule adapters (containers, db, auth, memory)
— agentic/	# Agentic behavior framework
— agents/	# CLI agent registry (48 agents)
— analytics/	# Usage analytics
— audit/	# Audit logging

	auth/	# JWT, OAuth credential reading
	background/	# Background task management
	benchmark/	# Performance benchmarks
	bigdata/	# Infinite context, distributed
memory		
	browser/	# Playwright browser automation
	cache/	# Caching layer (Redis)
	challenges/	# Challenge framework
	config/	# Centralized env-var
configuration		
	database/	# PostgreSQL connectivity +
migrations		
	debate/	# Debate orchestration
	embeddings/	# Embedding providers (13)
	eventbus/	# Event-driven architecture
	features/	# Feature flags
	formatters/	# Code formatters (32+)
	handlers/	# HTTP handlers (~40 files)
	health/	# Health check system
	llm/	# LLM abstraction layer
	providers/	# 51+ provider implementations
	ai21/	
	anthropic/	
	anthropic_cu/	
	azure/	
	cerebras/	
	chutes/	
	claude/	
	cloudflare/	
	codestral/	
	cohere/	
	deepseek/	
	fireworks/	
	gemini/	
	githubmodels/	
	groq/	
	huggingface/	
	hyperbolic/	
	junie/	
	kilo/	
	kimi/	
	kimicode/	
	lmstudio/	
	mistral/	
	modal/	
	nia/	
	nlpcloud/	
	novita/	
	nvidia/	
	ollama/	
	openai/	
	openrouter/	
	perplexity/	
	publicai/	

```

├── qwen/
├── replicate/
├── sambanova/
├── sarvam/
├── siliconflow/
├── together/
├── upstage/
├── venice/
├── vertex/
├── vulavula/
├── xai/
├── zai/
├── zen/
├── zhipu/
├── mcp/ # MCP server registry +
connection pooling
├── middleware/ # Auth, compression, concurrency
limits
├── models/ # Core domain types
├── ports/ # Centralized port registry
├── router/ # Gin router setup + service
initialization
├── security/ # Guardrails, red-team, PII
detection
├── services/ # Business logic layer
├── llmsverifier_score_adapter.go # LLMsVerifier score
adapter
├── provider_registry.go # Provider registry
├── debate_service.go # Debate orchestration
├── acp_client.go # ACP client
├── acp_manager.go # ACP manager
├── ...
├── streaming/ # Real-time streaming responses
├── transport/ # HTTP/3 (QUIC) + Brotli
compression
├── utils/ # Shared utilities
├── verifier/ # LLMsVerifier integration layer
├── doc.go # Package documentation
├── config.go # Verifier configuration
├── config_test.go # Config tests
├── discovery.go # Model discovery service
├── discovery_test.go # Discovery tests
├── database.go # SQLite database operations
├── database_test.go # Database tests
├── enhanced_scoring.go # 7-component scoring engine
├── health.go # Health monitoring
├── provider_types.go # UnifiedProvider, UnifiedModel
types
├── scoring.go # 5-component scoring engine
├── startup.go # Startup verification
orchestrator
├── verification.go # Verification service
├── adapters/
├── provider_adapter.go # Provider interface

```

```

adapter
|
|   └─ free_adapter.go           # Free provider
(Zen, OpenRouter)
|
|   └─ oauth_adapter.go         # OAuth provider
(Claude, Qwen)
|
|   └─ extended_providers_adapter.go # Extended
providers (Grok, Perplexity, etc.)
|
|   └─ extended_registry.go      # Extended provider
registry
|
|   └─ pkg/api/                  # Generated protobuf code
|   └─ tests/                    # Test suites organized by type
|   |   └─ unit/                 # Unit tests (mocks allowed,
-short)
|   |   └─ integration/          # Cross-service integration
tests
|   |   └─ e2e/                  # End-to-end workflows
|   |   └─ security/             # Vulnerability scans
|   |   └─ stress/               # Load/saturation tests
|   |   └─ chaos/                # Fault injection
|   |   └─ challenge/            # Competition tests
|   |   └─ performance/          # Benchmarks (//go:build
performance)
|   |   └─ fixtures/             # Shared test data
|   |   └─ testutils/            # Shared helpers
|
|   └─ challenges/               # 193+ validation scripts
|       └─ scripts/
|           └─ llmsverifier_cliagents_challenge.sh
|           └─ llmsverifier_startup_verification_challenge.sh
|           └─ llmsverifier_submodule_smoke_challenge.sh
|           └─ startup_verifier_debate_team_challenge.sh
|           └─ verifier_filtering_challenge.sh
|           └─ ... (190+ more)
|
|   └─ configs/
|       └─ verifier.yaml         # LLMsVerifier configuration
|
|   └─ docs/                     # Comprehensive documentation
|       └─ guides/llms-verifier.md
|       └─ integration/LLMSVERIFIER_INTEGRATION_PLAN.md
|       └─ verifier/
|           └─ API.md
|           └─ README.md
|           └─ USER_GUIDE.md
|           └─ LLMSVERIFIER_POWER_FEATURES.md
|       └─ ...
|
|   └─ LLMsVerifier/             # Git submodule (vasic-digital/
LLMsVerifier)
|       └─ llm-verifier/         # The actual verifier code
|           └─ api_keys/         # API key tracking
|           └─ pkg/cliagents/    # CLI agent unified generator
|           └─ ...

```

— MCP/	# MCP submodule with 35+
submodules	
— cli_agents/	# 48 CLI agent submodules
— Toolkit/	# Provider toolkits
— k8s/	# Kubernetes manifests
— monitoring/	# Prometheus + Grafana configs
— scripts/	# 90+ build/test/deploy scripts
— .gitmodules	# 60+ submodules defined
— go.mod	# Go module with replace
directives	
— package.json	# Playwright dependency only
— Makefile	# Comprehensive build/test
targets	
— Dockerfile	
— docker-compose.yml	
— .env.example	# Environment configuration
template	
— .env.mcp.example	# MCP configuration template
— CLAUDE.md	# AI agent hard stops + rules
— AGENTS.md	# Authoritative agent guide
— CONSTITUTION.md	# 33 mandatory rules
— README.md	

C. LLMsVerifier INTEGRATION ANALYSIS (MOST CRITICAL)

C.1 LLMsVerifier as Git Submodule

Location: LLMsVerifier/ (external submodule) **Repository:** `git@github.com:vasic-digital/LLMsVerifier.git` **Commit:** `1d53ae3b72c77c1f27171c0677431c48d2d02bdd`

C.2 Go Module Integration

go.mod dependency:

`digital.vasic.llmsverifier v0.0.0`

go.mod replace directive (maps local submodule):

// Part of the extensive replace block for local submodules

C.3 Import Statements (EXACT)

Import 1 - CLI agent configuration generation:

```
// File: cmd/helixagent/main.go
digital.vasic.llmsverifier/pkg/cliagents
```

Import 2 - API key management:

```
// File: internal/verifier/startup.go
digital.vasic.llmsverifier/api_keys
```

C.4 Initialization Pattern

File: cmd/helixagent/main.go

The LLMsVerifier is initialized as part of the StartupVerifier which orchestrates the complete provider verification pipeline:

```
// In main.go - the startup flow:
// 1. Load Config & Environment
// 2. Initialize StartupVerifier (Scoring + Verification + Health)
// 3. Discover ALL Providers (API Key + OAuth + Free)
// 4. Verify ALL Providers in Parallel (8-test pipeline)
// 5. Score ALL Verified Providers (5-component weighted)
// 6. Rank by Score
// 7. Select AI Debate Team (up to 25 LLMs)
// 8. Start Server with Verified Configuration
```

Key initialization in internal/verifier/startup.go:

```
// NewStartupVerifier creates the orchestrator
func NewStartupVerifier(cfg *StartupConfig, log *logrus.Logger)
    *StartupVerifier {
    // Initialize services
    verifierCfg := DefaultConfig()
    verifierSvc := NewVerificationService(verifierCfg)
    scoringSvc, err := NewScoringService(verifierCfg)

    // Initialize enhanced scoring service (Phase 1: 7-component
    // scoring)
    enhancedScoring := NewEnhancedScoringService(scoringSvc)

    return &StartupVerifier{
        config:          cfg,
        verifierSvc:      verifierSvc,
        scoringSvc:       scoringSvc,
        enhancedScoring: enhancedScoring,
        subscriptionDetector: NewSubscriptionDetector(log),
        // ...
    }
}
```

C.5 Configuration Routing (How API Keys Flow)

Configuration file: configs/verifier.yaml

```

verifier:
  enabled: true
  verification:
    mandatory_code_check: true
    code_visibility_prompt: "Do you see my code?"
    verification_timeout: 60s
    retry_count: 3
    retry_delay: 5s
    tests:
      - existence
      - responsiveness
      - latency
      - streaming
      - function_calling
      - coding_capability
      - error_detection
      - code_visibility
  scoring:
    weights:
      response_speed: 0.25
      model_efficiency: 0.20
      cost_effectiveness: 0.25
      capability: 0.20
      recency: 0.10
    cache_ttl: 24h

```

Environment Variables (from `.env.example`):

- OPENAI_API_KEY=sk-your-openai-key
- ANTHROPIC_API_KEY=sk-ant-your-anthropic-key
- DEEPSEEK_API_KEY=sk-your-deepseek-key
- GROQ_API_KEY=gsk-your-groq-key
- MISTRAL_API_KEY=your-mistral-key
- COHERE_API_KEY=your-cohere-key
- PERPLEXITY_API_KEY=pplx-your-perplexity-key
- GEMINI_API_KEY=your-gemini-key
- TOGETHER_API_KEY=your-together-key
- XAI_API_KEY=your-xai-key
- CEREBRAS_API_KEY=your-cerebras-key
- CLOUDFLARE_API_KEY=your-cloudflare-key
- SILICONFLOW_API_KEY=your-siliconflow-key

C.6 Model Consumption Pattern

5-Component Weighted Scoring:

Component	Weight	Description
ResponseSpeed	25%	API response latency
ModelEfficiency	20%	Token efficiency
CostEffectiveness	25%	Cost per token
Capability	20%	Model capability score
Recency	10%	Model release date

7-Component Enhanced Scoring (Phase 1):


Component	Weight	Description
ResponseSpeed	20%	API latency
ModelEfficiency	15%	Token efficiency
CostEffectiveness	20%	Cost per 1K

tokens || Capability | 15% | Model capability tier || Recency | 5% | Model release date || CodeQuality | 15% | Code generation benchmarks || ReasoningScore | 10% | Reasoning task performance |



C.7 All Files Referencing LLMsVerifier


- 1. cmd/helixagent/main.go - Imports digital.vasic.llmsverifier/pkg/cliagents
- 2. internal/verifier/startup.go - Imports digital.vasic.llmsverifier/api_keys
- 3. internal/services/llmsverifier_score_adapter.go - Score adapter connecting ProviderDiscovery to LLMsVerifier scoring
- 4. internal/services/llmsverifier_score_adapter_test.go - Tests for score adapter
- 5. configs/verifier.yaml - Configuration file
- 6. internal/verifier/ - Complete integration layer (~15 files)
- 7. challenges/scripts/llmsverifier_*.sh - 4 challenge scripts
- 8. docs/guides/llms-verifier.md
- 9. docs/integration/LLMSVERIFIER_INTEGRATION_PLAN.md
- 10. docs/verifier/*.md
- 11. challenge-results/llmsverifier-*/

C.8 Integration Architecture Diagram (Text-Based)









D. CONFIGURATION SYSTEM

D.1 All .env Variables

Key environment variable categories from `.env.example`:

LLM Provider API Keys:

```
OPENAI_API_KEY=sk-your-openai-key  
ANTHROPIC_API_KEY=sk-ant-your-anthropic-key
```

DEEPSEEK_API_KEY=sk-your-deepseek-key
GROQ_API_KEY=gsk-your-groq-key
MISTRAL_API_KEY=your-mistral-key
COHERE_API_KEY=your-cohere-key
PERPLEXITY_API_KEY=pplx-your-perplexity-key
GEMINI_API_KEY=your-gemini-key
TOGETHER_API_KEY=your-together-key
XAI_API_KEY=your-xai-key
CEREBRAS_API_KEY=your-cerebras-key
CLOUDFLARE_API_KEY=your-cloudflare-key
SILICONFLOW_API_KEY=your-siliconflow-key

Infrastructure:

DB_HOST=localhost
DB_PORT=8101
DB_USER=helixagent
DB_PASSWORD=helixagent123
REDIS_HOST=localhost
REDIS_PORT=8102
REDIS_PASSWORD=helixagent123
PORT=8100
GIN_MODE=release
JWT_SECRET=your-secret-key-here

HelixMemory:

HELIX_MEMORY_MODE=local
HELIX_MEMORY_COGNEE_ENDPOINT=http://localhost:8000
HELIX_MEMORY_MEM0_ENDPOINT=http://localhost:8001
HELIX_MEMORY_LETTA_ENDPOINT=http://localhost:8283

HelixLLM:

USE_HELIX_LLM=true
HELIX_LLM_ENDPOINT=http://localhost:8443
HELIX_LLM_MODE=full

Port Registry:

HELIXAGENT_PORT_PREFIX=8
HELIXAGENT_PORT_HTTP=8100
HELIXAGENT_PORT_POSTGRES=8101
HELIXAGENT_PORT_REDIS=8102
HELIXAGENT_PORT_MCP_BRIDGE=8103

D.2 How LLMsVerifier Gets Its Config

1. **Primary:** configs/verifier.yaml - Structured YAML with provider configs, scoring weights, test configuration
2. **Secondary:** Environment variables injected into YAML via \${VAR} interpolation
3. **Tertiary:** .env file loaded via godotenv at startup

D.3 All Providers Configured

From `configs/verifier.yaml`:

- **openai**: gpt-4, gpt-4-turbo, gpt-4o, gpt-3.5-turbo
- **anthropic**: claude-3-5-sonnet, claude-3-opus, claude-3-sonnet, claude-3-haiku
- **google**: gemini-1.5-pro, gemini-1.5-flash, gemini-pro
- **groq**: llama-3.3-70b, llama-3.1-8b, mixtral-8x7b
- **together**: meta-llama/Llama-3-70b, mistralai/Mixtral-8x7B
- **mistral**: mistral-large, mistral-medium, mistral-small
- **deepseek**: deepseek-chat, deepseek-coder
- **xai**: grok-beta (disabled by default)
- **cerebras**: llama-3.3-70b (disabled by default)
- **cloudflare**: @cf/meta/llama-2-7b (disabled by default)
- **siliconflow**: deepseek-ai/DeepSeek-V2.5 (disabled by default)
- **replicate**: meta/llama-2-70b (disabled by default)
- **openrouter**: anthropic/claude-3.5, openai/gpt-4, meta/llama-3.1-405b
- **ollama**: DEPRECATED (score: 5.0, development only)

D.4 MCP Configuration

From `.env.mcp.example`:

- 35 MCP server submodules configured
 - SSE bridge on port 8103
 - Auto-start on HelixAgent boot
 - Working-only mode (skip API-key-required MCPs by default)
-

E. ALL CAPABILITIES

E.1 Providers Supported (51+)

Tier 1 (Primary - Always enabled):

1. Claude (Anthropic)
2. DeepSeek
3. Gemini (Google)
4. Groq
5. Mistral
6. OpenAI
7. Together AI
8. OpenRouter
9. Qwen (Alibaba)

Tier 2 (Extended): 10. AI21 11. Anthropic (CU variant) 12. Azure 13. Cerebras 14. Chutes 15. Cloudflare Workers AI 16. Codestral 17. Cohere 18. Fireworks 19. GitHub Models 20. HuggingFace 21. Hyperbolic 22. Junie 23. Kilo 24. Kimi 25. Kimi Code 26. LM Studio 27. Modal 28. NIA 29. NLP Cloud 30. Novita 31. NVIDIA 32. Ollama (DEPRECATED) 33. Perplexity 34. PublicAI 35. Replicate 36. SambaNova 37. Sarvam 38. SiliconFlow 39. Upstage 40. Venice 41. Vertex (Google) 42. VulaVula 43. xAI (Grok) 44. ZAI 45. Zen (OpenCode) 46. Zhipu

E.2 MCPs Integrated (35+)

Full MCP submodule list from `.gitmodules`:

- microsoft-mcp, python-sdk, typescript-sdk
- brave-search, notion-mcp-server
- sentry-mcp, github-mcp-server
- playwright-mcp, browserbase-mcp
- qdrant-mcp, supabase-mcp, redis-mcp
- elasticsearch-mcp, obsidian-mcp, firecrawl-mcp
- cloudflare-mcp, workers-mcp, aws-mcp
- kubernetes-mcp, k8s-mcp-server
- slack-mcp, telegram-mcp
- airtable-mcp, trello-mcp, heroku-mcp
- mongodb-mcp, atlassian-mcp
- perplexity-mcp, omniseach-mcp
- context7-mcp, langchain-mcp, llamaindex-mcp
- docs-mcp, all-in-one-mcp

E.3 LSPs Integrated (10)

Referenced in documentation (server implementations via `sourcegraph/jsonrpc2`):

- gopls, pylsp, typescript-language-server
- rust-analyzer, clangd
- And 4 more

E.4 ACPs Integrated

- ACP Manager on port 8300
- JSON-RPC protocol support
- Tool calling + context management
- Agent registration/execution endpoints

E.5 Embeddings Systems (13)

Via `internal/embeddings/` + `Embeddings/` submodule:

- OpenAI, Cohere, Mistral embeddings
- Local embeddings (all-mpnet-base-v2)
- Provider-specific embedding endpoints

E.6 RAG Systems

- RAG/ submodule
- NVIDIA RAG container support
- ChromaDB, Qdrant, Neo4j vector stores
- Chunking: 1000 tokens, 200 overlap, `top_k=5`

E.7 Skills System

- Skill definitions in provider configurations
- Capability detection per model
- Specialization tagging (code, reasoning, general)

E.8 Plugins System

- Plugins/ submodule
 - Hot reloading support
 - Auto-discovery on startup
-

F. TEST & CHALLENGE STRUCTURE

F.1 Challenge Scripts (193+)

Located in challenges/scripts/. Key LLMsVerifier challenges:

1. **llmsverifier_cliagents_challenge.sh** - Tests unified CLI agent config generation
2. **llmsverifier_startup_verification_challenge.sh** - Tests API key tracking + provider verification
3. **llmsverifier_submodule_smoke_challenge.sh** - Tests submodule compilation
4. **startup_verifier_debate_team_challenge.sh** - Tests debate team selection
5. **verifier_filtering_challenge.sh** - Tests provider filtering by score

Challenge structure:

```
#!/bin/bash
# Color-coded output (GREEN=pass, RED=fail, YELLOW=test,
#                     BLUE=header)
# Manual pass/fail counting (TESTS_PASSED, TESTS_FAILED)
# Each test function: print_test → test logic → pass_test/
#                   fail_test
# Final summary with exit 0 (all pass) or exit 1 (any fail)
```

F.2 Test Files

Unit tests (mocks allowed, -short flag):

- internal/verifier/config_test.go - 11,346 bytes
- internal/verifier/database_test.go - 6,291 bytes
- internal/verifier/discovery_test.go - 18,344 bytes
- internal/services/llmsverifier_score_adapter_test.go
- All provider *_test.go files

Integration tests (REAL infrastructure required):

- tests/integration/ - Cross-service tests
- tests/e2e/ - End-to-end workflows
- tests/security/ - Vulnerability scans
- tests/stress/ - Load tests
- tests/chaos/ - Fault injection

F.3 Anti-Bluff Testing Patterns

From CLAUDE.md / CONSTITUTION.md:

- **CONST-002a**: NO mocks/stubs in production code
- **CONST-002**: 100% test coverage across ALL test types
- **CONST-030**: Real infrastructure for ALL non-unit tests
- **Rule**: "Mocks/stubs ONLY in unit tests; all other tests use real data and live services"
- **Enforcement**: make no-mocks-above-unit with strict allowlist ratchet

F.4 Test Frameworks Used

- **Go testing**: Standard go test + testify
 - **Playwright**: Browser automation tests
 - **Challenge scripts**: Bash-based validation (193+ scripts)
 - **HelixQA**: Quality assurance framework (separate submodule)
-

G. DOCUMENTATION GOVERNANCE

G.1 CLAUDE.md Analysis

File: CLAUDE.md (comprehensive - ~20,000+ bytes)

Hard Stops (NON-NEGOTIABLE):

1. **NO CI/CD pipelines** - No .github/workflows/, .gitlab-ci.yml, etc.
2. **NO manual container commands** - Only ./bin/helixagent orchestrates containers
3. **NO HTTPS for Git** - SSH URLs only (git@github.com:...)
4. **Run go mod vendor after touching submodules**

Constitution Rules Extracted:

- CONST-001 through CONST-033 (33 mandatory rules)
- CONST-029: Concurrent-safe containers (safe.Store/safe.Slice)
- CONST-030: Real infrastructure for all non-unit tests

G.2 AGENTS.md Analysis

File: AGENTS.md (authoritative guide for AI agents)

Key sections:

- Project overview (Go 1.26, module dev.helix.agent)
- Technology stack table
- Code organization diagram
- Build/test commands
- Code style rules
- Provider patterns

G.3 CONSTITUTION.md

File: CONSTITUTION.md

33 mandatory rules across 17 categories: | Category | Count | Key Rules |
|-----|-----|-----| | Quality | 2 | No broken components, no dead code | |
Safety | 1 | Memory safety | | Security | 1 | Security scanning | | Performance | 2 |
Monitoring, lazy loading | | Containerization | 4 | Full containerization, orchestration
flow | | Configuration | 2 | Unified config, non-interactive execution | | Testing | 8 |
100% coverage, challenges, stress tests | | Documentation | 2 | Complete docs,
synchronization | | Principles | 2 | KISS/DRY/SOLID, design patterns | | Stability | 1 |
Rock-solid changes | | Observability | 1 | Health and monitoring | | GitOps | 2 |
GitSpec compliance, SSH only | | CI/CD | 1 | Manual CI/CD only | | Networking | 1 |
HTTP/3 (QUIC) + Brotli | | Resource Management | 1 | Test resource limits | |
Concurrency | 1 | safe.Store/safe.Slice |

G.4 Quality Requirements

1. **100% Test Coverage** - Every component: unit, integration, E2E, security, stress, chaos, automation, benchmark
 2. **Challenge Coverage** - Every component **MUST** have challenge scripts
 3. **Containerization** - All services in containers
 4. **Real Data** - Beyond unit tests: actual API calls, real databases, live services
 5. **No False Positives** - Tests validate actual behavior, not return codes
-

H. DEPENDENCIES & PACKAGE.JSON

H.1 package.json

```
{  
  "dependencies": {  
    "playwright": "^1.58.2"  
  }  
}
```

Note: Playwright is used for browser automation testing only. The project is **Go-based**, not Node.js.

H.2 go.mod Key Dependencies

HelixAgent module: `dev.helix.agent` **Go version:** 1.26

Key external dependencies:

- `github.com/gin-gonic/gin v1.12.0` - HTTP framework
- `github.com/jackc/pgx/v5` - PostgreSQL driver
- `github.com/redis/go-redis/v9` - Redis client
- `github.com/sirupsen/logrus` - Structured logging
- `github.com/prometheus/client_golang` - Metrics
- `github.com/docker/docker` - Docker API
- `github.com/neo4j/neo4j-go-driver/v5` - Neo4j
- `github.com/ClickHouse/clickhouse-go/v2` - ClickHouse
- `github.com/minio/minio-go/v7` - MinIO/S3
- `github.com/playwright-community/playwright-go` - Playwright Go bindings
- `github.com/quic-go/quic-go` - HTTP/3 (QUIC)
- `github.com/andybalholm/brotli` - Brotli compression
- `go.opentelemetry.io/otel` - OpenTelemetry
- `google.golang.org/grpc` - gRPC

Internal submodules (41 extracted modules):

- `digital.vasic.llmsverifier` - LLMsVerifier
 - `digital.vasic.mcp` - MCP module
 - `digital.vasic.helixmemory` - Memory system
 - `digital.vasic.helixqa` - QA framework
 - `digital.vasic.helixspecifier` - Spec generation
 - `digital.vasic.concurrency` - Safe concurrent data structures
 - `digital.vasic.llmprovider` - LLM provider abstraction
 - `digital.vasic.llmorchestrator` - LLM orchestration
 - And 33 more...
-

I. KEY SOURCE FILES ANALYSIS

I.1 `internal/verifier/startup.go` (CRITICAL)

Purpose: Orchestrates the complete startup verification pipeline

Key functions:

- `NewStartupVerifier()` - Creates the orchestrator with all services
- `VerifyAllProviders()` - 5-phase verification pipeline
- `discoverProviders()` - Discovers API key + OAuth + free providers
- `verifyProviders()` - 8-test parallel verification
- `scoreProviders()` - 5/7-component weighted scoring
- `selectDebateTeam()` - Selects up to 25 LLMs for debate

LLMsVerifier imports used:

- `digital.vasic.llmsverifier/api_keys` - API key tracking

I.2 internal/services/llmsverifier_score_adapter.go (CRITICAL)

Purpose: Bridges ProviderDiscovery to LLMsVerifier scoring system

Key features:

- LLMsVerifierScoreAdapter struct with safe.Store for thread safety
- GetProviderScore() / GetModelScore() - Normalizes 0-100 to 0-10
- RefreshScores() - Periodic score refresh (5-minute interval)
- inferProviderFromModel() - Dynamic provider inference from model ID patterns

I.3 internal/verifier/discovery.go (CRITICAL)

Purpose: Automatic model discovery from all configured providers

Key features:

- ModelDiscoveryService - Concurrent-safe discovery
- discoverAllModels() - Parallel discovery from all provider endpoints
- verifyDiscoveredModels() - Parallel model verification
- selectTopModels() - Score-based model selection for ensemble

I.4 internal/verifier/scoring.go (CRITICAL)

Purpose: 5-component weighted scoring engine

Key features:

- ScoringService with cache (safe.Store)
- CalculateScore() - Dynamic model class inference
- inferModelClassScore() - Pattern-based scoring (not hardcoded)
- BatchCalculateScores() - Parallel score computation

I.5 internal/verifier/enhanced_scoring.go (CRITICAL)

Purpose: 7-component enhanced scoring for debate team selection

Key features:

- EnhancedScoringService with safe.Store cache
- CalculateEnhancedScore() - Full 7-component calculation
- calculateComponents() - All component score calculations
- calculateDiversityBonus() - Diversity-aware team selection

I.6 internal/services/provider_registry.go (CRITICAL)

Purpose: Manages LLM provider registration with LLMsVerifier integration

Key features:

- ProviderRegistry with safe.Store for all collections
- scoreAdapter field - LLMsVerifier score adapter
- startupVerifier field - Atomic pointer to startup verifier
- 51+ provider imports

I.7 internal/verifier/adapters/oauth_adapter.go

Purpose: Handles OAuth-based providers (Claude, Qwen)

Key insight: Claude OAuth tokens from Claude Code CLI are **PRODUCT-RESTRICTED** - they can only be used with Claude Code itself, not the standard Anthropic API. The adapter marks these as unverified and excludes them from the debate team.


I.8 internal/verifier/adapters/free_adapter.go


Purpose: Handles free providers (Zen/OpenCode, OpenRouter :free)

Key features:

- Two-phase verification (direct API → CLI facade fallback)
- Base score 6.0-7.0 for free providers
- ZenCLIProvider for models that fail direct API verification

J. INTEGRATION ARCHITECTURE SUMMARY





|



4. SCORE ADAPTER (`llmsverifier_score_adapter.go`)

- Normalizes scores ($0-100 \rightarrow 0-10$)
- Caches in `safe.Store`
- Refreshes every 5 minutes

|



5. PROVIDER REGISTRY (`provider_registry.go`)

- Uses scores for dynamic provider ordering
- Circuit breakers + concurrency limits

|



6. ENSEMBLE SERVICE → Routes requests to best providers

|



7. DEBATE SERVICE → Multi-round AI debates

- Uses debate team from startup verification
- 5 positions x 5 LLMs = 25 total participants

K. BLUEPRINT FOR HelixTranslate INTEGRATION

Based on this comprehensive analysis, the key integration points for adding LLMsVerifier to HelixTranslate are:

K.1 Required Components

1. **Git Submodule:** Add LLMsVerifier as a submodule (or copy key packages)
2. **Go Module Integration:** Add `digital.vasic.llmsverifier` to `go.mod` with `replace` directive
3. **Verifier Package:** Create `internal/verifier/` with:
 - `startup.go` - Startup verification orchestrator
 - `scoring.go` - 5-component scoring engine
 - `discovery.go` - Model discovery service
 - `provider_types.go` - Unified types
 - `adapters/` - Provider-specific adapters
4. **Score Adapter:** Create `internal/services/llmsverifier_score_adapter.go`
5. **Configuration:** Add `configs/verifier.yaml`

K.2 Critical Import Paths

```
// For CLI agent config generation
import "digital.vasic.llmsverifier/pkg/cliagents"

// For API key tracking
import "digital.vasic.llmsverifier/api_keys"
```

K.3 Configuration Pattern

```
# configs/verifier.yaml
verifier:
  enabled: true
  verification:
    mandatory_code_check: true
    tests: [existence, responsiveness, latency, streaming,
            code_visibility]
  scoring:
    weights:
      response_speed: 0.25
      model_efficiency: 0.20
      cost_effectiveness: 0.25
      capability: 0.20
      recency: 0.10
```

K.4 Startup Verification Pattern

```
// Create verifier
verifier := verifier.NewStartupVerifier(config, logger)
```

```
// Run verification
results, err := verifier.VerifyAllProviders(ctx)
if err != nil {
    log.Fatal("Verification failed:", err)
}

// Get debate team
team := verifier.SelectDebateTeam()
```

End of Comprehensive Analysis